

```

In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import SplineTransformer, OneHotEncoder
from sklearn.preprocessing import StandardScaler

#1. Load the dataset
housing = pd.read_csv("housing.csv")

#2. Create a stratified test set
housing["income_cat"] = pd.cut(housing['median_income'],
                               bins=[0, 1.5, 3.0, 4.5, 6.0, np.inf],
                               labels=[1, 2, 3, 4, 5])

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)

for train_index, test_index in split.split(housing, housing["income_cat"]):
    strat_train_set = housing.loc[train_index].drop("income_cat", axis=1)
    strat_test_set = housing.loc[test_index].drop("income_cat", axis=1)

# We will work on the copy of training data
housing = strat_train_set.copy()

#3. Seperate fetures and Labels
housing_labels = housing["median_house_value"].copy()
housing = housing.drop("median_house_value", axis=1)

print(housing, housing_labels)

#4. List and Seperate numerical and categorical columns
num_attribs = housing.drop("ocean_proximity", axis=1).columns.tolist()
cat_attribs = ["ocean_proximity"]

#5. Lets make the pipeline

# For numerical columns
num_pipeline = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("Standardize", StandardScaler())
])

# For Categoriocl columns
cat_pipeline = Pipeline([
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

# Construct the full pipeline
full_pipeline = ColumnTransformer([
    ("num", num_pipeline, num_attribs),
    ("cat", cat_pipeline, cat_attribs)
])

```

#6. Transform the data

```
housing_prepared = full_pipeline.fit_transform(housing)
print(housing_prepared)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
12655	-121.46	38.52	29	3873	797.0	
15502	-117.23	33.09	7	5320	855.0	
2908	-119.04	35.37	44	1618	310.0	
14053	-117.13	32.75	24	1877	519.0	
20496	-118.70	34.28	27	3536	646.0	
...	
15174	-117.07	33.03	14	6665	1231.0	
12661	-121.42	38.51	15	7901	1422.0	
19263	-122.72	38.44	48	707	166.0	
19140	-122.70	38.31	14	3155	580.0	
19773	-122.14	39.97	27	1079	222.0	

	population	households	median_income	ocean_proximity
12655	2237	706	2.1736	INLAND
15502	2015	768	6.3373	NEAR OCEAN
2908	667	300	2.8750	INLAND
14053	898	483	2.2264	NEAR OCEAN
20496	1837	580	4.4964	<1H OCEAN
...
15174	2026	1001	5.0900	<1H OCEAN
12661	4769	1418	2.8139	INLAND
19263	458	172	3.1797	<1H OCEAN
19140	1208	501	4.1964	<1H OCEAN
19773	625	197	3.1319	INLAND

[16512 rows x 9 columns] 12655 72100

15502 279600

2908 82700

14053 112500

20496 238300

...

15174 268500

12661 90400

19263 140400

19140 258100

19773 62700

Name: median_house_value, Length: 16512, dtype: int64

```
[[-0.94135046  1.34743822  0.02756357 ...  0.      0.
  0.      ]
 [ 1.17178212 -1.19243966 -1.72201763 ...  0.      0.
  1.      ]
 [ 0.26758118 -0.1259716   1.22045984 ...  0.      0.
  0.      ]
 ...
 [-1.5707942   1.31001828  1.53856552 ...  0.      0.
  0.      ]
 [-1.56080303  1.2492109   -1.1653327   ...  0.      0.
  0.      ]
 [-1.28105026  2.02567448 -0.13148926 ...  0.      0.
  0.      ]]
```

In [2]: `from sklearn.linear_model import LinearRegression`
`from sklearn.tree import DecisionTreeRegressor`

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import root_mean_squared_error
```

7. Train the model

Linear Regression

```
In [3]: # lin_reg = LinearRegression()
# lin_reg.fit(housing_prepared, housing_labels)
# lin_preds = lin_reg.predict(housing_prepared)
# # lin_rmse = root_mean_squared_error(housing_labels, lin_preds)
```

```
In [4]: # print(f"The root mean squared error for Linear Regression is {lin_rmse}")
```

Decision Tree

```
In [5]: #dec_reg = DecisionTreeRegressor()
#dec_reg.fit(housing_prepared, housing_labels)
#dec_preds = dec_reg.predict(housing_prepared)
#dec_rmse = root_mean_squared_error(housing_labels, dec_preds)
```

```
In [6]: # print(f"The root mean squared error for Desision tree is {dec_rmse}")
```

Random Forest Model

```
In [7]: # random_forest_reg = RandomForestRegressor()
# random_forest_reg.fit(housing_prepared, housing_labels)
# random_forest_preds = random_forest_reg.predict(housing_prepared)
# random_forest_rmse = root_mean_squared_error(housing_labels, random_forest_pred
```

```
In [8]: # print(f"The root mean squared error for Random Forest is {random_forest_rmse}")
```

8. Cross Validattion

```
In [9]: from sklearn.model_selection import cross_val_score
```

```
In [10]: dec_reg = DecisionTreeRegressor()
dec_reg.fit(housing_prepared, housing_labels)
dec_preds = dec_reg.predict(housing_prepared)
#dec_rmse = root_mean_squared_error(housing_labels, dec_preds)
dec_rmses = -cross_val_score(dec_reg, housing_prepared, housing_labels, scoring="neg_mean_squared_error")

print(pd.Series(dec_rmses).describe())

random_forest_reg = RandomForestRegressor()
random_forest_reg.fit(housing_prepared, housing_labels)
random_forest_preds = random_forest_reg.predict(housing_prepared)
# random_forest_rmse = root_mean_squared_error(housing_labels, random_forest_pred
random_forest_rmses = -cross_val_score(random_forest_reg, housing_prepared, housin

print(pd.Series(random_forest_rmses).describe())
```

```
lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)
lin_preds = lin_reg.predict(housing_prepared)
#lin_rmse = root_mean_squared_error(housing_labels, lin_preds)
lin_rmses = -cross_val_score(lin_reg, housing_prepared, housing_labels, scoring="neg_mean_squared_error", cv=5)

print(pd.Series(lin_rmses).describe())
```

```
count      10.000000
mean      69329.806956
std       2717.228989
min       63732.461319
25%      68030.605234
50%      69593.605180
75%      71136.802180
max       73163.033739
dtype: float64
count      10.000000
mean      49370.008603
std       2089.715241
min       46048.804428
25%      47851.130605
50%      49023.234196
75%      50631.029659
max       52990.201716
dtype: float64
count      10.000000
mean      69204.322755
std       2500.382157
min       65318.224029
25%      67124.346106
50%      69404.658178
75%      70697.800632
max       73003.752739
dtype: float64
```

In []:

In []:

In []:

In []: